

## Assignment- Logistic Regression (Theoretical)

### #Q1. What is Logistic Regression, and how does it differ from Linear Regression?

Logistic Regression is a classification algorithm that predicts the probability of an instance belonging to a particular class.

It differs from Linear Regression as it outputs probabilities (bounded between 0 and 1) instead of continuous values.

### #Q2. What is the mathematical equation of Logistic Regression?

The mathematical equation of Logistic Regression is based on the sigmoid function, which maps any real-valued number to a probability between 0 and 1.

$$\bullet \quad P(Y = 1|X) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 X_1 + \dots + \beta_n X_n)}}$$

### #Q3. Why do we use the Sigmoid function in Logistic Regression

The **Sigmoid function** (also called the **logistic function**) is used in Logistic Regression because it converts any real-valued input into a probability between **0 and 1**. This makes it ideal for **binary classification problems**, where we need to determine whether an instance belongs to class **1 (positive)** or **0 (negative)**.

### #Q4. What is the cost function of Logistic Regression

In Logistic Regression, we use the **Log Loss (Logarithmic Loss)** or **Binary Cross-Entropy** as the cost function instead of Mean Squared Error (MSE) used in Linear Regression. This is because MSE is not suitable for classification problems, as it results in a **non-convex** function, making optimization difficult.

### #Q5. What is Regularization in Logistic Regression?

**Regularization** is a technique used to **prevent overfitting** by adding a penalty term to the cost function. It discourages large coefficients, making the model more **generalizable**.

### #Q6 Why is it needed Explain the difference between Lasso, Ridge, and Elastic Net regression

#### Types of Regularization:

1. **Lasso (L1 Regularization)** – Shrinks coefficients and forces some to become **zero** (feature selection).
2. **Ridge (L2 Regularization)** – Shrinks coefficients but **does not** force them to zero.
3. **Elastic Net (L1 + L2)** – A combination of **Lasso and Ridge**, balancing feature selection and coefficient shrinkage.

### #Q7. When should we use Elastic Net instead of Lasso or Ridge

Elastic Net is preferred when:

- Features are **highly correlated**: Lasso may randomly select one feature and ignore others, but Elastic Net keeps a balance.
- Data is **high-dimensional** (many more features than samples): It combines L1 (Lasso) and L2 (Ridge) penalties to provide stability.
- You need **automatic feature selection** (Lasso-like effect) while preventing **over-shrinking** of coefficients (Ridge-like effect).

### #Q8. What is the impact of the regularization parameter ( $\lambda$ ) in Logistic Regression

The regularization parameter  $\lambda$  (or  $C = \frac{1}{\lambda}$  in scikit-learn) controls the complexity of the model:

- **Higher  $\lambda$**  → Stronger penalty → Simpler model (less overfitting) → Smaller coefficients.
- **Lower  $\lambda$**  → Weaker penalty → More flexible model (but risk of overfitting) → Larger coefficients.

### #Q9. What are the key assumptions of Logistic Regression?

1. **Linear Relationship (with Logit)**: The independent variables should have a linear relationship with the log-odds.
2. **No Multicollinearity**: Features should not be highly correlated; otherwise, it affects coefficient stability.
3. **Independent Observations**: Data points should be independent of each other.
4. **Large Sample Size**: Logistic Regression performs best when there are enough samples, especially for rare classes.

### #Q10. What are some alternatives to Logistic Regression for classification tasks?

- **Decision Trees** (handles non-linearity well).
- **Random Forests** (ensemble learning, reduces overfitting).
- **Support Vector Machines (SVMs)** (works well with high-dimensional data).
- **Naive Bayes** (good for text classification and probabilistic modeling).
- **Neural Networks** (for complex, non-linear relationships).
- **Gradient Boosting (XGBoost, LightGBM, CatBoost)** (for structured data).

### #Q11. What are Classification Evaluation Metrics?

- **Accuracy** =  $(TP + TN) / (TP + TN + FP + FN)$
- **Precision** =  $TP / (TP + FP)$  (how many predicted positives are correct).
- **Recall (Sensitivity)** =  $TP / (TP + FN)$  (how many actual positives are identified).
- **F1-score** =  $2 \times (Precision \times Recall) / (Precision + Recall)$  (harmonic mean of precision and recall).
- **ROC-AUC Score** (discrimination between classes).
- **Log Loss** (penalizes wrong probability estimates).

### #Q12. How does class imbalance affect Logistic Regression?

- Logistic Regression may **favor the majority class**, predicting the minority class poorly.
- Metrics like **accuracy** become misleading.
- **Solutions:**
  - **Resampling techniques** (oversampling minority or undersampling majority class).
  - **Use class weights** (`class_weight='balanced'` in sklearn).
  - **Use different metrics** like Precision-Recall AUC instead of accuracy.

### #Q13. What is Hyperparameter Tuning in Logistic Regression?

Hyperparameter tuning optimizes model performance. Key parameters:

- **Regularization parameter ( $\lambda$  or C):** Controls overfitting.
- **Penalty type (L1, L2, Elastic Net):** Chooses the right regularization method.
- **Solver selection** (affects optimization performance).
- **Grid Search / Random Search:** Common tuning methods.

### #Q14. What are different solvers in Logistic Regression? Which one should be used?

- **liblinear:** Good for small datasets, supports L1 & L2.
- **lbfgs:** Default for small to medium datasets, supports only L2.
- **saga:** Good for large datasets, supports L1, L2, and Elastic Net.
- **newton-cg & cg:** Works well for only L2 regularization.

Which one to use?

- For **small datasets** → `liblinear`.
- For **large datasets** → `saga`.
- For **L2 regularization** → `lbfgs` or `newton-cg`.

#### #Q15. How is Logistic Regression extended for multiclass classification?

- **One-vs-Rest (OvR):** Fits separate binary classifiers for each class.
- **Softmax (Multinomial Regression):** Assigns probabilities to multiple classes simultaneously.
- **Which one to use?**
  - **OvR:** Works well with fewer classes, interpretable.
  - **Softmax:** Preferred when classes are mutually exclusive.

#### #Q16. What are the advantages and disadvantages of Logistic Regression?

##### Advantages:

- Simple, interpretable, and efficient.
- Works well with linearly separable data.
- Probabilistic output (confidence scores).
- Less prone to overfitting with proper regularization.

##### Disadvantages:

- Struggles with non-linear relationships.
- Sensitive to outliers.
- Requires feature scaling for best performance.

#### #Q17. What are some use cases of Logistic Regression?

- **Healthcare:** Disease prediction (e.g., diabetes, cancer detection).
- **Finance:** Credit risk analysis, fraud detection.
- **Marketing:** Customer churn prediction.
- **Human Resources:** Employee attrition prediction.
- **Political Science:** Predicting election outcomes.

#### #Q18. What is the difference between Softmax Regression and Logistic Regression?

- **Logistic Regression:** Used for **binary classification**, predicts probability for two classes.
- **Softmax Regression:** Generalization of logistic regression for **multiclass classification**, assigns probability to each class.

#### #Q19. How do we choose between One-vs-Rest (OvR) and Softmax for multiclass classification?

**Use OvR** when:

- The number of classes is large.
- Faster training time.

**Use Softmax** when:

- Classes are mutually exclusive.
- Probabilistic interpretation is needed.

### **#Q20. How do we interpret coefficients in Logistic Regression?**

Each coefficient  $\beta_i$  represents the change in log-odds for a one-unit increase in the corresponding feature.

$e^{\beta_i}$  gives the odds ratio (how much more/less likely an event is).

If  $\beta_i > 0$ , the feature increases the probability of the positive class.

If  $\beta_i < 0$ , the feature decreases the probability of the positive class.